

# ML Tutorial 2

June 13, 2018

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import sklearn
import seaborn as sns
sns.set()
%matplotlib inline
```

## 0.1 Arrays & Dataframes

Recall we use arrays (and array of arrays) to deal with data in numpy:

```
In [2]: data = np.random.multivariate_normal([155,65],[[5,3.5],[3.5,5]],20)
print(data)
```

```
[[153.81830494  66.75216825]
 [153.24958843  63.97775387]
 [155.53157868  66.35131989]
 [155.74792032  65.20016683]
 [157.0232017   68.94099499]
 [154.81196077  62.67780174]
 [155.32994637  64.57465533]
 [153.73971327  66.58747603]
 [153.39280592  66.46613757]
 [157.75488641  66.94391405]
 [157.19444109  65.92912883]
 [156.33356568  64.09789433]
 [158.1872879   68.0998284 ]
 [149.7580223   57.92745011]
 [155.14979759  65.49127875]
 [156.25311415  67.28307103]
 [157.6431965   68.52903368]
 [154.30779624  63.09863841]
 [156.99808757  64.89770124]
 [153.49262644  63.87615423]]
```

We can reshape the array using the `.reshape` method

```
In [3]: print(data.reshape(-1,4))

[[153.81830494  66.75216825 153.24958843  63.97775387]
 [155.53157868  66.35131989 155.74792032  65.20016683]
 [157.0232017   68.94099499 154.81196077  62.67780174]
 [155.32994637  64.57465533 153.73971327  66.58747603]
 [153.39280592  66.46613757 157.75488641  66.94391405]
 [157.19444109  65.92912883 156.33356568  64.09789433]
 [158.1872879   68.0998284  149.7580223   57.92745011]
 [155.14979759  65.49127875 156.25311415  67.28307103]
 [157.6431965   68.52903368 154.30779624  63.09863841]
 [156.99808757  64.89770124 153.49262644  63.87615423]]
```

Note that when reshaping, you can use the `-1` flag to indicate an automatic calculation of the column/row>

Sometimes dealing with arrays is a pain, using dataframes is a way to solve the problem:

```
In [4]: df = pd.DataFrame(data,columns=("h","w"))

pd.DataFrame.head(df)
```

```
Out [4]:
```

	h	w
0	153.818305	66.752168
1	153.249588	63.977754
2	155.531579	66.351320
3	155.747920	65.200167
4	157.023202	68.940995

We can add a column by simply assigning a new column into the dataframe, this immediately implies that dataframe is **mutable** :

```
In [5]: df['r'] = np.sqrt((df['w'] + df['h'])*2)

pd.DataFrame.head(df)
```

```
Out [5]:
```

	h	w	r
0	153.818305	66.752168	220.570473
1	153.249588	63.977754	217.227342
2	155.531579	66.351320	221.882899
3	155.747920	65.200167	220.948087
4	157.023202	68.940995	225.964197

You can delete columns using the `.drop` method, however this method does not cause the original dataframe to mutate, unless you have the `inplace` parameter set to `True`:

```
In [6]: df.drop(columns=['r'])

pd.DataFrame.head(df)
```

```
Out [6]:
```

	h	w	r
0	153.818305	66.752168	220.570473
1	153.249588	63.977754	217.227342
2	155.531579	66.351320	221.882899
3	155.747920	65.200167	220.948087
4	157.023202	68.940995	225.964197

```
In [7]: df.drop(columns=['r'],inplace=True)

pd.DataFrame.head(df)

df['r'] = np.sqrt((df['w'] + df['h'])**2)
```

Most of the time you DO NOT want to delete data. A commonly used best practice is to simply create another data frame for the data you are trying to manipulate:

```
In [8]: df2 = df.drop(columns=['r'])

pd.DataFrame.head(df2)
```

```
Out [8]:
```

	h	w
0	153.818305	66.752168
1	153.249588	63.977754
2	155.531579	66.351320
3	155.747920	65.200167
4	157.023202	68.940995

We can merge dataframes with the same column names together:

```
In [9]: D = pd.DataFrame(np.arange(0,40).reshape(-1,2),columns=("A","B"))
pd.DataFrame.head(D)
```

```
Out [9]:
```

	A	B
0	0	1
1	2	3
2	4	5
3	6	7
4	8	9

```
In [10]: E = pd.DataFrame(np.arange(41,51).reshape(-1,2),columns=("A","B"))
pd.DataFrame.head(E)
```

```
Out [10]:
```

	A	B
0	41	42
1	43	44
2	45	46
3	47	48
4	49	50

```
In [11]: combined_df = pd.concat([D,E])
combined_df
```

```
Out[11]:
```

	A	B
0	0	1
1	2	3
2	4	5
3	6	7
4	8	9
5	10	11
6	12	13
7	14	15
8	16	17
9	18	19
10	20	21
11	22	23
12	24	25
13	26	27
14	28	29
15	30	31
16	32	33
17	34	35
18	36	37
19	38	39
0	41	42
1	43	44
2	45	46
3	47	48
4	49	50

If you try to concatenate two dataframes with different column indices, pandas will encourage you to do the concatenation along the columns this results in the concatenation based on the `iloc`:

```
In [12]: F = pd.DataFrame(np.arange(0,40).reshape(-1,2),columns=("A1","B1"))
```

```
G = pd.DataFrame(np.arange(41,51).reshape(-1,2),columns=("A2","B2"))
```

```
combined_df2 = pd.concat([F,G],axis = 1)
combined_df2
```

```
Out[12]:
```

	A1	B1	A2	B2
0	0	1	41.0	42.0
1	2	3	43.0	44.0
2	4	5	45.0	46.0
3	6	7	47.0	48.0
4	8	9	49.0	50.0
5	10	11	NaN	NaN
6	12	13	NaN	NaN
7	14	15	NaN	NaN

8	16	17	NaN	NaN
9	18	19	NaN	NaN
10	20	21	NaN	NaN
11	22	23	NaN	NaN
12	24	25	NaN	NaN
13	26	27	NaN	NaN
14	28	29	NaN	NaN
15	30	31	NaN	NaN
16	32	33	NaN	NaN
17	34	35	NaN	NaN
18	36	37	NaN	NaN
19	38	39	NaN	NaN

If you attempt to concatenate the two dataframes along the rows, pandas will return the following:

```
In [13]: combined_df3 = pd.concat([F,G])
         combined_df3
```

```
/usr/local/lib/python3.6/site-packages/ipykernel_launcher.py:1: FutureWarning: Sorting because
of pandas will change to not sort by default.
```

To accept the future behavior, pass 'sort=True'.

To retain the current behavior and silence the warning, pass sort=False

```
"""Entry point for launching an IPython kernel.
```

```
Out[13]:
```

	A1	A2	B1	B2
0	0.0	NaN	1.0	NaN
1	2.0	NaN	3.0	NaN
2	4.0	NaN	5.0	NaN
3	6.0	NaN	7.0	NaN
4	8.0	NaN	9.0	NaN
5	10.0	NaN	11.0	NaN
6	12.0	NaN	13.0	NaN
7	14.0	NaN	15.0	NaN
8	16.0	NaN	17.0	NaN
9	18.0	NaN	19.0	NaN
10	20.0	NaN	21.0	NaN
11	22.0	NaN	23.0	NaN
12	24.0	NaN	25.0	NaN
13	26.0	NaN	27.0	NaN
14	28.0	NaN	29.0	NaN
15	30.0	NaN	31.0	NaN
16	32.0	NaN	33.0	NaN
17	34.0	NaN	35.0	NaN
18	36.0	NaN	37.0	NaN

19	38.0	NaN	39.0	NaN
0	NaN	41.0	NaN	42.0
1	NaN	43.0	NaN	44.0
2	NaN	45.0	NaN	46.0
3	NaN	47.0	NaN	48.0
4	NaN	49.0	NaN	50.0

This results in a dataframe containing complementary information from the two prior dataframes.

## 0.2 Plotting using Matplotlib and Seaborn

Let's consider the iris dataset.

```
In [14]: iris = sns.load_dataset("iris")
```

```
In [15]: iris.head()
```

```
Out[15]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

Let's find out the number of unique species in this dataset:

```
In [16]: iris['species'].unique()
```

```
Out[16]: array(['setosa', 'versicolor', 'virginica'], dtype=object)
```


Now let's explore the data set by plotting a few graphs:

### 0.2.1 Scatter Plots

If we want to visualise two continuous variables, we would try to view it as a **scatter plot**. Let's try to plot the petal length against petal width:

```
In [17]: plt.scatter(iris['petal_length'],iris['petal_width'])
```

```
Out[17]: <matplotlib.collections.PathCollection at 0x10f7a7ba8>
```



## 0.2.2 Histograms

If we want to visualise a single continuous variable, we would want to try to view its distribution over a number line. Let's try plotting a **histogram** to observe the distribution of the petal length and petal widths, we set the number of bins (bars) to 20:


```
In [18]: plt.hist(iris['petal_length'],bins=20)
```

```
Out[18]: (array([ 4., 33., 11.,  2.,  0.,  0.,  1.,  2.,  3.,  5., 12., 14., 12.,
        17.,  6., 12.,  7.,  4.,  2.,  3.]),
         array([1.    , 1.295, 1.59 , 1.885, 2.18 , 2.475, 2.77 , 3.065, 3.36 ,
        3.655, 3.95 , 4.245, 4.54 , 4.835, 5.13 , 5.425, 5.72 , 6.015,
        6.31 , 6.605, 6.9  ]),
         <a list of 20 Patch objects>)
```



```
In [19]: plt.hist(iris['petal_width'],bins=20)
```

```
Out[19]: (array([34.,  7.,  7.,  1.,  1.,  0.,  0.,  7.,  3.,  5., 21., 12.,  4.,
                2., 12., 11.,  6.,  3.,  8.,  6.]),
          array([0.1 , 0.22, 0.34, 0.46, 0.58, 0.7 , 0.82, 0.94, 1.06, 1.18, 1.3 ,
                1.42, 1.54, 1.66, 1.78, 1.9 , 2.02, 2.14, 2.26, 2.38, 2.5 ]),
          <a list of 20 Patch objects>)
```






### 0.2.3 Joint Plots

Seaborn allows us to plot visualise the prior plots together. Let's now visualise the two variables together in one plot:

```
In [20]: sns.jointplot(iris['petal_width'],iris['petal_length'],kind = 'scatter')
```

```
/usr/local/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning: The 'normed'
  warnings.warn("The 'normed' kwarg is deprecated, and has been "
```

```
Out[20]: <seaborn.axisgrid.JointGrid at 0x10f916c18>
```




## 0.2.4 Box Plots

Supposing now that we want to compare the distributions for each species of flowers, one of the most common ways to present this information is through a box plot. In general this is a good idea if you want to compare continuous variable against a categorical variable.

Let's compare the Sepal Length for each Species:

```
In [21]: sns.boxplot(iris['species'],iris['sepal_length'])
```


```
Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x10ff4e7b8>
```



And again for the Petal Length:

```
In [22]: sns.boxplot(iris['species'],iris['petal_length'])
```

```
Out[22]: <matplotlib.axes._subplots.AxesSubplot at 0x110080d68>
```



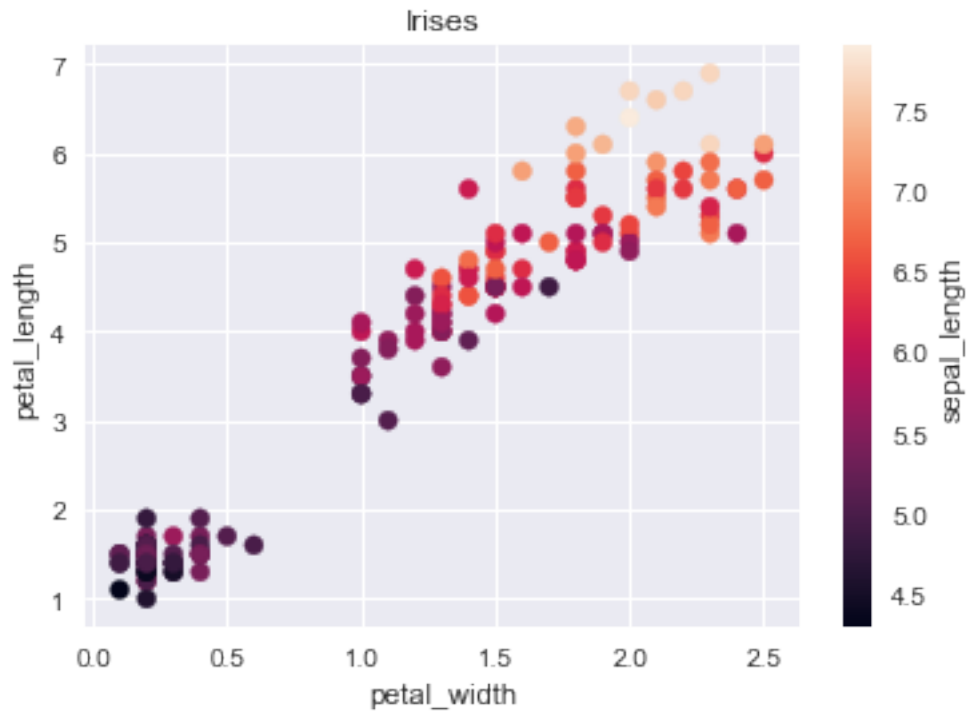
### 0.3 Scatter Plots with added Dimensions

What if we want to consider three variables at the same time? We can always segment by a new dimension using size or colour:

```
In [23]: plt.scatter(iris['petal_width'],iris['petal_length'],c=iris['sepal_length'])

plt.colorbar().set_label('sepal_length')
plt.xlabel('petal_width')
plt.ylabel('petal_length')
plt.title("Irises")
```

```
Out[23]: Text(0.5,1,'Irises')
```

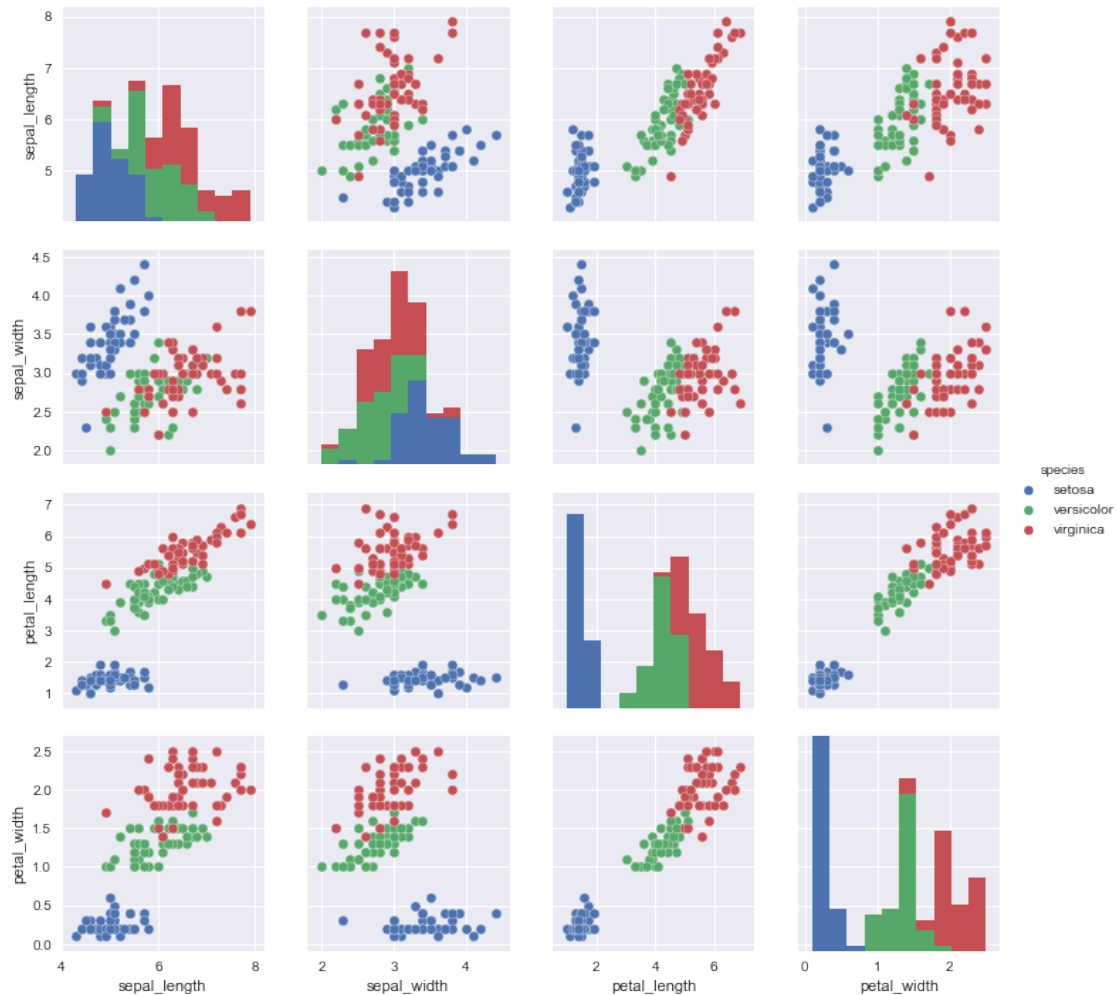


## 0.4 Pairwise Plotting

In general, we would always try to plot the pairs to observe all the interactions:

```
In [24]: sns.pairplot(iris,hue="species")
```

```
Out[24]: <seaborn.axisgrid.PairGrid at 0x110260be0>
```



## 0.5 Linear Regression

We now consider a new dataset. First, get the data from 442 diabetes patients and the progression of their disease from the scikit-learn library.

```
In [25]: from sklearn import datasets
         from sklearn.linear_model import LinearRegression
```

```
In [26]: diabetes = datasets.load_diabetes()
```

Let us view the description of the dataset.

```
In [27]: print(diabetes.DESCR)
```

```
Diabetes dataset
=====
```

## Notes

-----

Ten baseline variables, age, sex, body mass index, average blood pressure, and six blood serum measurements were obtained for each of  $n = 442$  diabetes patients, as well as the response of interest, a quantitative measure of disease progression one year after baseline.

### Data Set Characteristics:

:Number of Instances: 442

:Number of Attributes: First 10 columns are numeric predictive values

:Target: Column 11 is a quantitative measure of disease progression one year after baseline

:Attributes:

:Age:

:Sex:

:Body mass index:

:Average blood pressure:

:S1:

:S2:

:S3:

:S4:

:S5:

:S6:

Note: Each of these 10 feature variables have been mean centered and scaled by the standard deviation

Source URL:

<http://www4.stat.ncsu.edu/~boos/var.select/diabetes.html>

For more information see:

Bradley Efron, Trevor Hastie, Iain Johnstone and Robert Tibshirani (2004) "Least Angle Regression" ([http://web.stanford.edu/~hastie/Papers/LARS/LeastAngle\\_2002.pdf](http://web.stanford.edu/~hastie/Papers/LARS/LeastAngle_2002.pdf))

The features given by the dataset are as follows:

In [28]: `diabetes.feature_names`

Out[28]: `['age', 'sex', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6']`

We are aiming to predict the quantitative measure of disease progression one year after baseline (column 11), based on these features. We note that:

**Each of these 10 feature variables have been mean centered and scaled by the standard deviation times  $n_{\text{samples}}$  (i.e. the sum of squares of each column totals 1).**

## 0.5.1 Exploratory Data Analysis

Observing the first few rows of the dataset yields:

```
In [29]: npdata = diabetes.data
         features_df = pd.DataFrame(npdata, columns = diabetes.feature_names)
         target_df = pd.DataFrame(diabetes.target, columns = ['target'])
         diabetes_df = pd.concat([features_df, target_df], axis=1)
         pd.DataFrame.head(diabetes_df)
```

```
Out [29]:
```

	age	sex	bmi	bp	s1	s2	s3	\
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	
2	0.085299	0.050680	0.044451	-0.005671	-0.045599	-0.034194	-0.032356	
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038	
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142	

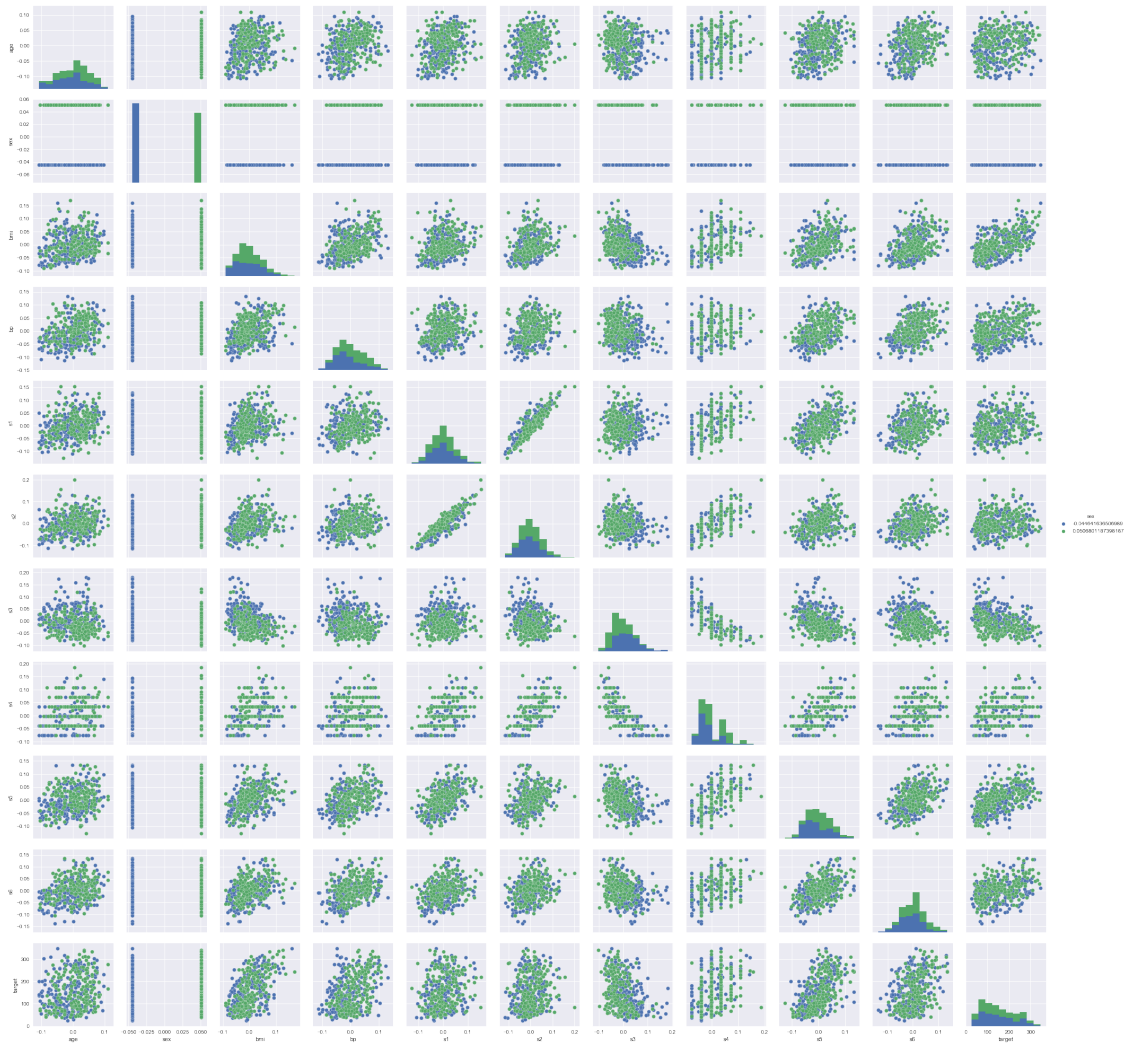
  

	s4	s5	s6	target
0	-0.002592	0.019908	-0.017646	151.0
1	-0.039493	-0.068330	-0.092204	75.0
2	-0.002592	0.002864	-0.025930	141.0
3	0.034309	0.022692	-0.009362	206.0
4	-0.002592	-0.031991	-0.046641	135.0

Let's try viewing the pairwise plots for each feature pair:

```
In [30]: sns.pairplot(diabetes_df, hue="sex")
```

```
Out [30]: <seaborn.axisgrid.PairGrid at 0x1119a9780>
```



```
In [31]: diabetes_df.corr()
```

```
Out[31]:
```

	age	sex	bmi	bp	s1	s2	s3 \
age	1.000000	0.173737	0.185085	0.335427	0.260061	0.219243	-0.075181
sex	0.173737	1.000000	0.088161	0.241013	0.035277	0.142637	-0.379090
bmi	0.185085	0.088161	1.000000	0.395415	0.249777	0.261170	-0.366811
bp	0.335427	0.241013	0.395415	1.000000	0.242470	0.185558	-0.178761
s1	0.260061	0.035277	0.249777	0.242470	1.000000	0.896663	0.051519
s2	0.219243	0.142637	0.261170	0.185558	0.896663	1.000000	-0.196455
s3	-0.075181	-0.379090	-0.366811	-0.178761	0.051519	-0.196455	1.000000
s4	0.203841	0.332115	0.413807	0.257653	0.542207	0.659817	-0.738493
s5	0.270777	0.149918	0.446159	0.393478	0.515501	0.318353	-0.398577
s6	0.301731	0.208133	0.388680	0.390429	0.325717	0.290600	-0.273697
target	0.187889	0.043062	0.586450	0.441484	0.212022	0.174054	-0.394789



	s4	s5	s6	target
age	0.203841	0.270777	0.301731	0.187889
sex	0.332115	0.149918	0.208133	0.043062
bmi	0.413807	0.446159	0.388680	0.586450
bp	0.257653	0.393478	0.390429	0.441484
s1	0.542207	0.515501	0.325717	0.212022
s2	0.659817	0.318353	0.290600	0.174054
s3	-0.738493	-0.398577	-0.273697	-0.394789
s4	1.000000	0.617857	0.417212	0.430453
s5	0.617857	1.000000	0.464670	0.565883
s6	0.417212	0.464670	1.000000	0.382483
target	0.430453	0.565883	0.382483	1.000000

On first glance, we can see a good sense of positive correlatedness between the target and bmi, and a negative correlatedness between target and s3. Further observation of the s1 and s2 variables shows that we can drop one of the variables from the analysis as they are extremely correlated (corrcoef = 0.9).

Now let's try fitting a linear model to the data.

Let's try the following model:

$$Target = \beta_0 + \beta_1 bmi + \epsilon$$

```
In [32]: Y = np.array(diabetes_df['target'])
        X = np.array(diabetes_df['bmi'])

        X = X.reshape(-1,1)
        print(X.shape)
```

```
(442, 1)
```

Now, fit the model using the LinearRegression function from sklearn:

```
In [33]: model = LinearRegression()
```

```
In [34]: model.fit(X,Y)
```

```
/usr/local/lib/python3.6/site-packages/sklearn/linear_model/base.py:509: RuntimeWarning: intern
linalg.lstsq(X, y)
```

```
Out [34]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

Checking the coefficient of determination gives us:

```
In [35]: model.coef_
```

```
Out [35]: array([949.43526038])
```

The intercept term is estimated as:

```
In [38]: model.intercept_
```

```
Out[38]: 152.1334841628967
```

We can then use the model to predict the **expected value** of the target variable for a particular value of  $bmi = 0.07$ :

```
In [37]: model.predict(0.07)
```

```
Out[37]: array([218.59395239])
```