

# ML Tutorial 5

June 25, 2018

## 0.1 models

Regression	Classification
LinearRegression	LogisticRegression ( <b>penalty</b> : 'l2', <b>solver</b> : {'liblinear'})
Lasso (L1 regularized) ( <b>alpha</b> : 1)	LogisticRegression ( <b>alpha</b> : 1, <b>penalty</b> : 'l1', <b>solver</b> : {'liblinear', 'sag'})
Ridge (L2 regularized) ( <b>alpha</b> : 1)	LogisticRegression ( <b>alpha</b> :1, <b>penalty</b> : 'l2', <b>solver</b> : {'liblinear', 'newton-cg', 'lbfgs', 'sag', 'saga'})
ElasticNet ( <b>alpha</b> : 1, <b>l1_ratio</b> : 0.5)	LogisticRegression ( <b>alpha</b> :1, <b>l1_ratio</b> : 0.5, <b>multi_class</b> : 'multinomial', <b>solver</b> : {'newton-cg', 'lbfgs', 'sag', 'saga'})
KNeighborsRegressor ( <b>n_neighbors</b> : 5)	KNeighborsClassifier ( <b>n_neighbors</b> : 5)
DecisionTreeRegressor ( <b>max_depth</b> : 3, <b>min_samples_leaf</b> : 10)	DecisionTreeClassifier ( <b>max_depth</b> : 3, <b>min_samples_leaf</b> : 10)
RandomForestRegressor ( <b>n_estimators</b> : 10, <b>max_depth</b> : 3, <b>min_samples_leaf</b> : 10)	RandomForestClassifier ( <b>n_estimators</b> : 10, <b>max_depth</b> : 3, <b>min_samples_leaf</b> : 10)
VotingClassifier ( <b>voting</b> : {'hard', 'soft'})	

## 0.2 metrics for n-fold cross-validation and test set performance

Regression	Classification
MSE (lower is better)	Classification Error (lower is better)
neg_mean_squared_error (higher is better)	accuracy (higher is better)
AUC	
f1	

## 0.3 Using `GridSearchCV`

```
from sklearn import datasets
from sklearn.linear_models import LogisticRegression
from sklearn.model_selection import GridSearchCV
```

```
iris = datasets.load_iris()
```

```

parameters = {'alpha':[0.1, 1, 10], 'solver':['liblinear', 'lbfgs', 'sag']}
log_clf = LogisticRegression(penalty='l2')
clf = GridSearchCV(log_clf, parameters, scoring="accuracy")
clf.fit(iris.data, iris.target)

```

```

>> sorted(clf.cv_results_.keys())
...
['mean_fit_time', 'mean_score_time', 'mean_test_score',...
 'mean_train_score', 'param_C', 'param_kernel', 'params',...
 'rank_test_score', 'split0_test_score',...
 'split0_train_score', 'split1_test_score', 'split1_train_score',...
 'split2_test_score', 'split2_train_score',...
 'std_fit_time', 'std_score_time', 'std_test_score', 'std_train_score'...]

# best parameter combination is the row which has the best score/rank of 1
>> pd.DataFrame(clf.cv_results_)

```

```

In [94]: %matplotlib inline

```

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import sklearn
import seaborn as sns
sns.set()

```

```

In [95]: iris.head()

```

```

Out[95]:   sepal_length  sepal_width  petal_length  petal_width  species
0          5.1           3.5           1.4           0.2    setosa
1          4.9           3.0           1.4           0.2    setosa
2          4.7           3.2           1.3           0.2    setosa
3          4.6           3.1           1.5           0.2    setosa
4          5.0           3.6           1.4           0.2    setosa

```

### 0.3.1 before starting on trees

```

# remember to run the following installs

```

```

> brew install graphviz
> pip3 install pydotplus

```

```

In [100]: from sklearn.tree import DecisionTreeClassifier

```

```

iris = sns.load_dataset('iris')

clf = DecisionTreeClassifier(max_depth=3)

x = iris.drop("species", axis=1)
y = iris["species"]
clf.fit(x,y)

```

```
Out[100]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=3,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                                splitter='best')
```

```
In [101]: x.columns
```

```
Out[101]: Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_width'], dtype='object')
```

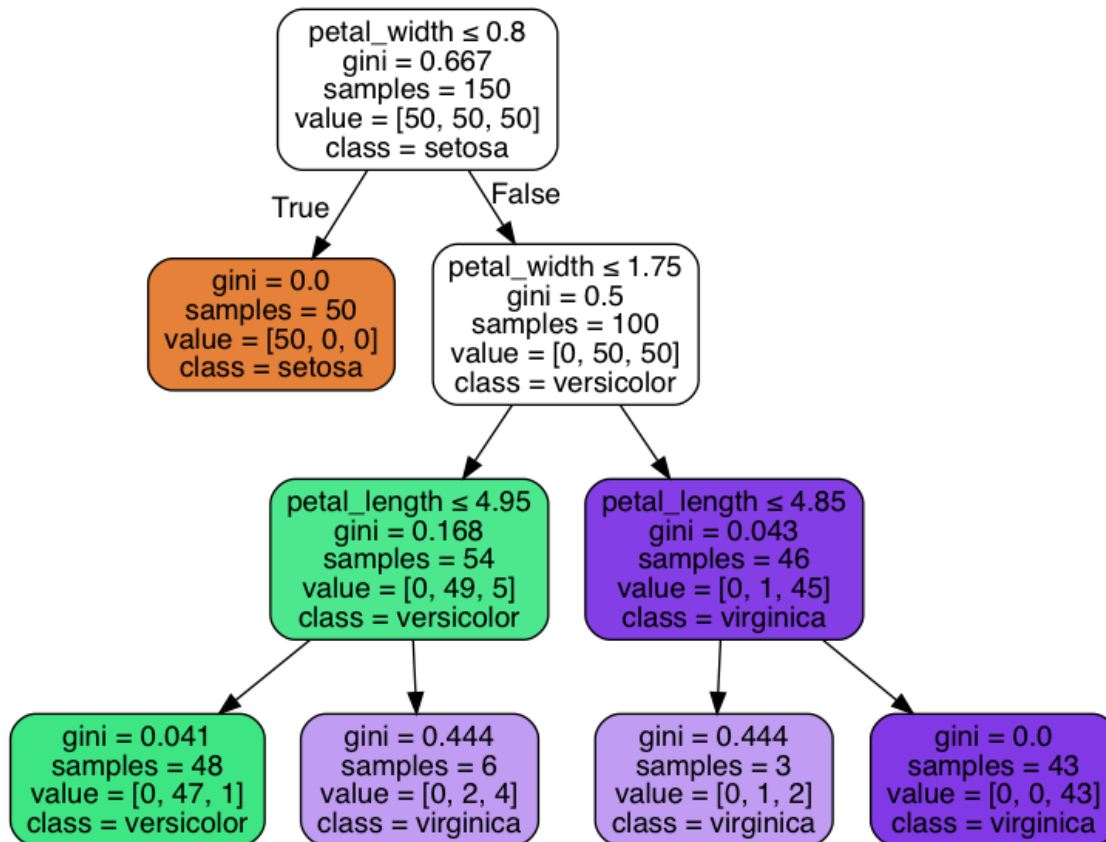
```
In [102]: y.unique()
```

```
Out[102]: array(['setosa', 'versicolor', 'virginica'], dtype=object)
```

```
In [103]: from sklearn.tree import export_graphviz
          from sklearn.externals.six import StringIO
          from IPython.display import Image
          import pydotplus

          dot_data = StringIO()
          export_graphviz(clf, out_file=dot_data,
                          filled=True,
                          rounded=True,
                          class_names=y.unique(),
                          feature_names=x.columns,
                          special_characters=True)
          graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
          Image(graph.create_png())
```

```
Out[103]:
```



```

In [ ]: from sklearn.datasets import load_iris
        from sklearn.tree import DecisionTreeClassifier

        # Parameters
        n_classes = 3
        plot_colors = "ryb"
        plot_step = 0.02

        # Load data
        iris = load_iris()

        for pairidx, pair in enumerate([[0, 1], [0, 2], [0, 3],
                                         [1, 2], [1, 3], [2, 3]]):
            # We only take the two corresponding features
            X = iris.data[:, pair]
            y = iris.target

            # Train
            clf = DecisionTreeClassifier().fit(X, y)

            # Plot the decision boundary

```

```

plt.subplot(2, 3, pairidx + 1)

x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, plot_step),
                     np.arange(y_min, y_max, plot_step))
plt.tight_layout(h_pad=0.5, w_pad=0.5, pad=2.5)

Z = clf.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)
cs = plt.contourf(xx, yy, Z, cmap=plt.cm.RdYlBu)

plt.xlabel(iris.feature_names[pair[0]])
plt.ylabel(iris.feature_names[pair[1]])

# Plot the training points
for i, color in zip(range(n_classes), plot_colors):
    idx = np.where(y == i)
    plt.scatter(X[idx, 0], X[idx, 1], c=color, label=iris.target_names[i],
                cmap=plt.cm.RdYlBu, edgecolor='black', s=15)

plt.suptitle("Decision surface of a decision tree using paired features")
plt.legend(loc='lower right', borderpad=0, handletextpad=0)
plt.axis("tight")
plt.show()

```

In [109]: `from scipy.special import comb`

```

def prob_majority_is_right(N, p):
    start = N//2 + 1
    total=0
    for k in range(start, N+1):
        total += comb(N,k) * p**k * (1-p)**(N-k)
    return total

for N in range(1,1001,10):
    print("P(majority is right for size {0}): \n {1:.2f}".format(N, prob_majority_is_

P(majority is right for size 1):
0.51
P(majority is right for size 11):
0.53
P(majority is right for size 21):
0.54
P(majority is right for size 31):
0.54
P(majority is right for size 41):
0.55

```

P(majority is right for size 51):  
0.56  
P(majority is right for size 61):  
0.56  
P(majority is right for size 71):  
0.57  
P(majority is right for size 81):  
0.57  
P(majority is right for size 91):  
0.58  
P(majority is right for size 101):  
0.58  
P(majority is right for size 111):  
0.58  
P(majority is right for size 121):  
0.59  
P(majority is right for size 131):  
0.59  
P(majority is right for size 141):  
0.59  
P(majority is right for size 151):  
0.60  
P(majority is right for size 161):  
0.60  
P(majority is right for size 171):  
0.60  
P(majority is right for size 181):  
0.61  
P(majority is right for size 191):  
0.61  
P(majority is right for size 201):  
0.61  
P(majority is right for size 211):  
0.61  
P(majority is right for size 221):  
0.62  
P(majority is right for size 231):  
0.62  
P(majority is right for size 241):  
0.62  
P(majority is right for size 251):  
0.62  
P(majority is right for size 261):  
0.63  
P(majority is right for size 271):  
0.63  
P(majority is right for size 281):  
0.63

P(majority is right for size 291):  
0.63  
P(majority is right for size 301):  
0.64  
P(majority is right for size 311):  
0.64  
P(majority is right for size 321):  
0.64  
P(majority is right for size 331):  
0.64  
P(majority is right for size 341):  
0.64  
P(majority is right for size 351):  
0.65  
P(majority is right for size 361):  
0.65  
P(majority is right for size 371):  
0.65  
P(majority is right for size 381):  
0.65  
P(majority is right for size 391):  
0.65  
P(majority is right for size 401):  
0.66  
P(majority is right for size 411):  
0.66  
P(majority is right for size 421):  
0.66  
P(majority is right for size 431):  
0.66  
P(majority is right for size 441):  
0.66  
P(majority is right for size 451):  
0.66  
P(majority is right for size 461):  
0.67  
P(majority is right for size 471):  
0.67  
P(majority is right for size 481):  
0.67  
P(majority is right for size 491):  
0.67  
P(majority is right for size 501):  
0.67  
P(majority is right for size 511):  
0.67  
P(majority is right for size 521):  
0.68

P(majority is right for size 531):  
0.68  
P(majority is right for size 541):  
0.68  
P(majority is right for size 551):  
0.68  
P(majority is right for size 561):  
0.68  
P(majority is right for size 571):  
0.68  
P(majority is right for size 581):  
0.69  
P(majority is right for size 591):  
0.69  
P(majority is right for size 601):  
0.69  
P(majority is right for size 611):  
0.69  
P(majority is right for size 621):  
0.69  
P(majority is right for size 631):  
0.69  
P(majority is right for size 641):  
0.69  
P(majority is right for size 651):  
0.70  
P(majority is right for size 661):  
0.70  
P(majority is right for size 671):  
0.70  
P(majority is right for size 681):  
0.70  
P(majority is right for size 691):  
0.70  
P(majority is right for size 701):  
0.70  
P(majority is right for size 711):  
0.70  
P(majority is right for size 721):  
0.70  
P(majority is right for size 731):  
0.71  
P(majority is right for size 741):  
0.71  
P(majority is right for size 751):  
0.71  
P(majority is right for size 761):  
0.71



P(majority is right for size 771):  
0.71  
P(majority is right for size 781):  
0.71  
P(majority is right for size 791):  
0.71  
P(majority is right for size 801):  
0.71  
P(majority is right for size 811):  
0.72  
P(majority is right for size 821):  
0.72  
P(majority is right for size 831):  
0.72  
P(majority is right for size 841):  
0.72  
P(majority is right for size 851):  
0.72  
P(majority is right for size 861):  
0.72  
P(majority is right for size 871):  
0.72  
P(majority is right for size 881):  
0.72  
P(majority is right for size 891):  
0.72  
P(majority is right for size 901):  
0.73  
P(majority is right for size 911):  
0.73  
P(majority is right for size 921):  
0.73  
P(majority is right for size 931):  
0.73  
P(majority is right for size 941):  
0.73  
P(majority is right for size 951):  
0.73  
P(majority is right for size 961):  
0.73  
P(majority is right for size 971):  
0.73  
P(majority is right for size 981):  
0.73  
P(majority is right for size 991):  
0.74

```
In [ ]: # VotingClassifier

from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import VotingClassifier

log_clf = LogisticRegression()
rnd_clf = RandomForestClassifier()
knn_clf = KNeighborsClassifier()

voting_clf = VotingClassifier(voting='hard',
                              estimators=[('lr', log_clf), ('rf', rnd_clf), ('knn', knn_clf)])
```